

Problem Statment :

Logistics and supply chain company which is the largest and fastest-growing fully integrated player in India by revenue in Fiscal 2021. They aim to build the operating system for commerce, through a combination of world-class infrastructure, logistics operations of the highest quality, and cutting-edge engineering and technology capabilities.

The Data team builds intelligence and capabilities using this data that helps them to widen the gap between the quality, efficiency, and profitability of their business versus their competitors.

Objective :

The company wants to understand and process the data coming out of data engineering pipelines:

- Clean, sanitize and manipulate data to get useful features out of raw fields
- Make sense out of the raw data and help the data science team to build forecasting models on it

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv('delhivery_data.csv')
```

```
In [ ]: # Let us check the data first.
df.head()
```

Out[]:

	data	trip_creation_time	route_schedule_uuid	route_type	trip_uuid
0	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	trip- 153741093647649320
1	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	trip- 153741093647649320
2	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	trip- 153741093647649320
3	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	trip- 153741093647649320
4	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	trip- 153741093647649320

5 rows × 24 columns



In []:

```
from google.colab import drive
drive.mount('/content/drive')
```

EDA

In []:

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 24 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   data                                144867 non-null  object
 1   trip_creation_time                 144867 non-null  object
 2   route_schedule_uuid               144867 non-null  object
 3   route_type                        144867 non-null  object
 4   trip_uuid                         144867 non-null  object
 5   source_center                    144867 non-null  object
 6   source_name                      144574 non-null  object
 7   destination_center               144867 non-null  object
 8   destination_name                 144606 non-null  object
 9   od_start_time                   144867 non-null  object
10  od_end_time                     144867 non-null  object
11  start_scan_to_end_scan           144867 non-null  float64
12  is_cutoff                       144867 non-null  bool
13  cutoff_factor                   144867 non-null  int64
14  cutoff_timestamp                 144867 non-null  object
15  actual_distance_to_destination    144867 non-null  float64
16  actual_time                     144867 non-null  float64
17  osrm_time                       144867 non-null  float64
18  osrm_distance                   144867 non-null  float64
19  factor                          144867 non-null  float64
20  segment_actual_time              144867 non-null  float64
21  segment_osrm_time               144867 non-null  float64
22  segment_osrm_distance            144867 non-null  float64
23  segment_factor                  144867 non-null  float64
dtypes: bool(1), float64(10), int64(1), object(12)
memory usage: 25.6+ MB

```

```

In [ ]: #We can see the numnber of rows and columns
print(f'Rows :{df.shape[0]}')
print(f'Columns :{df.shape[1]}')

```

```

Rows :144867
Columns :24

```

```

In [3]: # Droping the unwanted and Unknown field
df.drop(['is_cutoff', 'cutoff_factor', 'cutoff_timestamp', 'factor', 'segment_factor'])

```

```

In [ ]: #Let us check now if the cloumns are dropped
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   data                                  144867 non-null  object
1   trip_creation_time                   144867 non-null  object
2   route_schedule_uuid                 144867 non-null  object
3   route_type                           144867 non-null  object
4   trip_uuid                            144867 non-null  object
5   source_center                        144867 non-null  object
6   source_name                          144574 non-null  object
7   destination_center                   144867 non-null  object
8   destination_name                     144606 non-null  object
9   od_start_time                       144867 non-null  object
10  od_end_time                          144867 non-null  object
11  start_scan_to_end_scan               144867 non-null  float64
12  actual_distance_to_destination       144867 non-null  float64
13  actual_time                          144867 non-null  float64
14  osrm_time                           144867 non-null  float64
15  osrm_distance                        144867 non-null  float64
16  segment_actual_time                  144867 non-null  float64
17  segment_osrm_time                   144867 non-null  float64
18  segment_osrm_distance                144867 non-null  float64
dtypes: float64(8), object(11)
memory usage: 21.0+ MB

```

```

In [ ]: # Finding the null values
df.isnull().sum()
# We can see their are null vaulus in coulumn source_name and destination_name

```

Out[]:

	0
data	0
trip_creation_time	0
route_schedule_uuid	0
route_type	0
trip_uuid	0
source_center	0
source_name	293
destination_center	0
destination_name	261
od_start_time	0
od_end_time	0
start_scan_to_end_scan	0
actual_distance_to_destination	0
actual_time	0
osrm_time	0
osrm_distance	0
segment_actual_time	0
segment_osrm_time	0
segment_osrm_distance	0

dtype: int64

```
In [ ]: df[df['source_name'].isnull()]['source_center'].value_counts()
```

Out[]:


	count
source_center	
IND282002AAD	128
IND342902A1B	90
IND126116AAA	20
IND509103AAC	17
IND577116AAA	16
IND465333A1B	6
IND841301AAC	5
IND505326AAB	5
IND331022A1B	3
IND852118A1B	3

dtype: int64

In []: df[df['destination_name'].isnull()].head(3)

Out[]:

	data	trip_creation_time	route_schedule_uuid	route_type	trip_uui
110	training	2018-09-25 08:53:04.377810	thanos::sroute:4460a38d- ab9b-484e-bd4e- f4201d0...	FTL	15378655843775665
111	training	2018-09-25 08:53:04.377810	thanos::sroute:4460a38d- ab9b-484e-bd4e- f4201d0...	FTL	15378655843775665
982	test	2018-10-01 20:56:18.155260	thanos::sroute:d0ebdacd- e09b-47d3-be77- c9c4a05...	FTL	15384273781549566



- The above source centers do not have the name . Which can be difficult to track or of the drivers to know which place they would go .So it is better to maintain the name
- But for our analysis we can fill the null values with source centers

In []: df[df['destination_name'].isnull()]['destination_center'].value_counts()

Out[]:

destination_center	count
IND282002AAD	151
IND342902A1B	16
IND577116AAA	16
IND852118A1B	15
IND505326AAB	11
IND126116AAA	10
IND841301AAC	9
IND250002AAC	9
IND509103AAC	9
IND122015AAC	8
IND465333A1B	3
IND331001A1C	3
IND221005A1A	1

dtype: int64

```

In [ ]: destination_centers = [
    'IND282002AAD', 'IND342902A1B', 'IND577116AAA', 'IND852118A1B',
    'IND505326AAB', 'IND126116AAA', 'IND841301AAC', 'IND250002AAC',
    'IND509103AAC', 'IND122015AAC', 'IND465333A1B', 'IND331001A1C',
    'IND221005A1A'
]

for center in destination_centers:
    print(f"\nDestination Center: {center}")
    print(df[df['destination_center'] == center]['destination_name'].value_count)
# we can see evrery destination_name is null

```

Destination Center: IND282002AAD
 Series([], Name: count, dtype: int64)

Destination Center: IND342902A1B
 Series([], Name: count, dtype: int64)

Destination Center: IND577116AAA
 Series([], Name: count, dtype: int64)

Destination Center: IND852118A1B
 Series([], Name: count, dtype: int64)

Destination Center: IND505326AAB
 Series([], Name: count, dtype: int64)

Destination Center: IND126116AAA
 Series([], Name: count, dtype: int64)

Destination Center: IND841301AAC
 Series([], Name: count, dtype: int64)

Destination Center: IND250002AAC
 Series([], Name: count, dtype: int64)

Destination Center: IND509103AAC
 Series([], Name: count, dtype: int64)

Destination Center: IND122015AAC
 Series([], Name: count, dtype: int64)

Destination Center: IND465333A1B
 Series([], Name: count, dtype: int64)

Destination Center: IND331001A1C
 Series([], Name: count, dtype: int64)

Destination Center: IND221005A1A
 Series([], Name: count, dtype: int64)

- Same for above Destination Center they all have null Destination name .
- We can fill null as Destination Center in Destination name column.

In [4]: *# have fill the null values in the columns as their source center and Destination
 #in their respective name columns*

```
df['source_name']=df['source_name'].fillna(df['source_center'])
```

In [5]: `df['destination_name']=df['destination_name'].fillna(df['destination_center'])`

In []: *#Now Let us see the data type*
`df.info()`


```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 19 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   data                                144867 non-null  object
 1   trip_creation_time                  144867 non-null  object
 2   route_schedule_uuid                144867 non-null  object
 3   route_type                         144867 non-null  object
 4   trip_uuid                          144867 non-null  object
 5   source_center                      144867 non-null  object
 6   source_name                        144867 non-null  object
 7   destination_center                 144867 non-null  object
 8   destination_name                   144867 non-null  object
 9   od_start_time                      144867 non-null  object
10   od_end_time                        144867 non-null  object
11   start_scan_to_end_scan             144867 non-null  float64
12   actual_distance_to_destination     144867 non-null  float64
13   actual_time                        144867 non-null  float64
14   osrm_time                          144867 non-null  float64
15   osrm_distance                      144867 non-null  float64
16   segment_actual_time                144867 non-null  float64
17   segment_osrm_time                  144867 non-null  float64
18   segment_osrm_distance              144867 non-null  float64
dtypes: float64(8), object(11)
memory usage: 21.0+ MB

```

- we can see the columns which has time are object data type so we can change to datetime data type so that we can easily get the data from the column.

```
In [6]: df['trip_creation_time']=pd.to_datetime(df['trip_creation_time'],format='%Y-%m-%d %H:%M:%S.%f')
```


```
In [7]: df['od_start_time']=pd.to_datetime(df['od_start_time'],format='%Y-%m-%d %H:%M:%S.%f')
df['od_end_time']=pd.to_datetime(df['od_end_time'],format='%Y-%m-%d %H:%M:%S.%f')
```

```
In [ ]: df.duplicated().sum()
```

```
Out[ ]: np.int64(0)
```

- I want to visualize how many training and how many testing are their
- I also want to know about route_type

```
In [ ]: sns.countplot(x=df['data'])
plt.title('Frequency of data')
plt.show()
```




- As we can see there is more training data compared to test data .

```
In [ ]: #route_type

route_counts = df['route_type'].value_counts()
plt.pie(route_counts, labels=route_counts.index, autopct='%1.1f%%', startangle=90)
plt.title('Distribution of Route Types')
plt.show()
```

Distribution of Route Types



In []: `df.head(1)`

Out []:

	data	trip_creation_time	route_schedule_uuid	route_type	trip_uuid
0	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	trip- 153741093647649320

In [9]: `df['Source_state']=df['source_name'].str.split('(').str[1].str[:-1]`
`df['destination_state']=df['destination_name'].str.split('(').str[1].str[:-1]`

Feature Engineering and Hypothesis testing

In []: `df.groupby('trip_uuid')['trip_uuid'].value_counts()`

Out[]:

	count
trip_uuid	
trip-153671041653548748	39
trip-153671042288605164	9
trip-153671043369099517	89
trip-153671046011330457	2
trip-153671052974046625	7
...	...
trip-153861095625827784	7
trip-153861104386292051	2
trip-153861106442901555	6
trip-153861115439069069	17
trip-153861118270144424	4

14817 rows × 1 columns

dtype: int64In []: `df[df['trip_uuid']=='trip-153671052974046625']`

Out[]:

	data	trip_creation_time	route_schedule_uuid	route_type	trip_u
78217	training	2018-09-12 00:02:09.740725	thanos::sroute:d9f07b12- 65e0-4f3b-bec8- df06134...	FTL	153671052974046
78218	training	2018-09-12 00:02:09.740725	thanos::sroute:d9f07b12- 65e0-4f3b-bec8- df06134...	FTL	153671052974046
78219	training	2018-09-12 00:02:09.740725	thanos::sroute:d9f07b12- 65e0-4f3b-bec8- df06134...	FTL	153671052974046
78220	training	2018-09-12 00:02:09.740725	thanos::sroute:d9f07b12- 65e0-4f3b-bec8- df06134...	FTL	153671052974046
78221	training	2018-09-12 00:02:09.740725	thanos::sroute:d9f07b12- 65e0-4f3b-bec8- df06134...	FTL	153671052974046
78222	training	2018-09-12 00:02:09.740725	thanos::sroute:d9f07b12- 65e0-4f3b-bec8- df06134...	FTL	153671052974046
78223	training	2018-09-12 00:02:09.740725	thanos::sroute:d9f07b12- 65e0-4f3b-bec8- df06134...	FTL	153671052974046



In []:

```
df[df['trip_uuid']=='trip-153861118270144424']
```

Out[]:

	data	trip_creation_time	route_schedule_uuid	route_type	trip_uuid
11570	test	2018-10-03 23:59:42.701692	thanos::sroute:412fea14- 6d1f-4222-8a5f- a517042...	FTL	trip-153861118270144424
11571	test	2018-10-03 23:59:42.701692	thanos::sroute:412fea14- 6d1f-4222-8a5f- a517042...	FTL	trip-153861118270144424
11572	test	2018-10-03 23:59:42.701692	thanos::sroute:412fea14- 6d1f-4222-8a5f- a517042...	FTL	trip-153861118270144424
11573	test	2018-10-03 23:59:42.701692	thanos::sroute:412fea14- 6d1f-4222-8a5f- a517042...	FTL	trip-153861118270144424



In []:

```
# I can also drop segment wise time because we do not know the segment
#df.drop(['segment_actual_time', 'segment_osrm_time', 'segment_osrm_distance'], axis=1)
```

In [54]:

```
# Now grouping the the common trip_uuid and performing the agg function on it
data_centers=df.groupby(['trip_uuid', 'source_center', 'destination_center', 'data'])
```

```
'actual_dis
'actual_tim
'osrm_time'
'osrm_dista
}).reset_in
```

```
In [55]: data_centers.groupby(['source_center', 'destination_center'])['trip_uuid'].count
```

```
Out[55]:
```

		trip_uuid
source_center	destination_center	

IND562132AAA	IND560300AAA	151
	IND560099AAB	127
IND560099AAB	IND560300AAA	121
IND560300AAA	IND562132AAA	108
IND411033AAA	IND421302AAG	107

dtype: int64

- We can see the above most frequently used corridors
 1. IND562132AAA->IND560300AAA which is 151
 2. IND562132AAA ->IND560099AAB which is 127

```
In [10]: data=df.groupby(['trip_uuid', 'Source_state', 'destination_state', 'data']).agg({'o
'actual_dis
'actual_tim
'osrm_time'
'osrm_dista
'segment_act
'segment_osr
'segment_osr
}).reset_in
```

```
In [58]: data.describe()
```

Out[58]:

	od_start_time	od_end_time	start_scan_to_end_scan	actual_distance_to_
count	16657	16657	16657.000000	16657.000000
mean	2018-09-22 15:09:29.658076672	2018-09-22 21:32:38.367513344	401.975086	401.975086
min	2018-09-12 00:00:16.535741	2018-09-12 00:50:10.814399	23.000000	23.000000
25%	2018-09-17 06:08:40.561678080	2018-09-17 12:54:28.577978112	127.000000	127.000000
50%	2018-09-22 06:06:01.573842944	2018-09-22 13:33:44.977102080	219.000000	219.000000
75%	2018-09-27 19:31:17.301129984	2018-09-28 02:15:40.360954880	445.000000	445.000000
max	2018-10-06 04:27:23.392375	2018-10-08 03:00:24.353479	7898.000000	7898.000000
std	NaN	NaN	522.558609	522.558609



In [61]:

```
col=['actual_time','osrm_time','osrm_distance','actual_distance_to_destination']
for i in col:
    a=data[i].mean()
    print(f'Average of {i} :{a}')
```

Average of actual_time :273.54385543615297

Average of osrm_time :121.69952572492045

Average of osrm_distance :87.92251465658266

Average of actual_distance_to_destination :71.02438237948297

Average OSRM Distance: 87.92

Average OSRM Time: 121.69

Average Actual Time: 273.54385543615297

In []: data.shape

Out[]: (16657, 14)

- I have merged the rows with commn trip_uuid source_center,destination_center and took the numarical coloumn .
- Here we can observe for one trip is around trip like you can see below.

In []: data.head()

Out[]:

	trip_uuid	Source_state	destination_state	data	od_start_time	od_end_time
0	trip-153671041653548748	Madhya Pradesh	Uttar Pradesh	training	2018-09-12 00:00:16.535741	2018-09-12 16:39:46.858469
1	trip-153671041653548748	Uttar Pradesh	Haryana	training	2018-09-12 16:39:46.858469	2018-09-12 13:40:21.123456
2	trip-153671042288605164	Karnataka	Karnataka	training	2018-09-12 02:03:09.655591	2018-09-12 03:01:56.789012
3	trip-153671043369099517	Haryana	Punjab	training	2018-09-14 03:40:17.106733	2018-09-14 17:34:56.789012
4	trip-153671043369099517	Karnataka	Haryana	training	2018-09-12 00:00:33.691250	2018-09-12 03:40:17.106733



- Need to check the corellation

```
In [16]: data['Source_state'].value_counts()
# maximum soure_state is maharashtra
#Lower is Tripura
#We can take some action were other Lower service used states to use this Logist
```


Out[16]:

	count
Source_state	
Maharashtra	2797
Karnataka	2416
Haryana	1839
Tamil Nadu	1150
Uttar Pradesh	978
Telangana	838
Gujarat	819
Delhi	790
West Bengal	682
Punjab	663
Rajasthan	564
Andhra Pradesh	551
Madhya Pradesh	455
Bihar	398
Kerala	320
Assam	290
Jharkhand	189
Uttarakhand	182
Orissa	178
Chandigarh	160
Himachal Pradesh	143
Goa	74
Arunachal Pradesh	46
Chhattisgarh	43
Jammu & Kashmir	34
Pondicherry	19
Dadra and Nagar Haveli	15
Meghalaya	13
Mizoram	5
Nagaland	5
Tripura	1

dtype: int64

```
In [ ]: data['destination_state'].value_counts()  
#Maximum destination state is Maharashtra  
#Lower destination state is Daman & Diu
```

Out[]:

destination_state	count
Maharashtra	2705
Karnataka	2468
Haryana	1815
Tamil Nadu	1127
Uttar Pradesh	988
Telangana	902
Gujarat	819
Punjab	743
West Bengal	713
Delhi	674
Rajasthan	602
Andhra Pradesh	551
Madhya Pradesh	454
Bihar	403
Kerala	304
Assam	272
Jharkhand	197
Orissa	190
Uttarakhand	180
Himachal Pradesh	161
Chandigarh	123
Goa	74
Arunachal Pradesh	52
Chhattisgarh	43
Jammu & Kashmir	32
Pondicherry	24
Dadra and Nagar Haveli	17
Meghalaya	13
Mizoram	7
Tripura	2
Nagaland	1
Daman & Diu	1

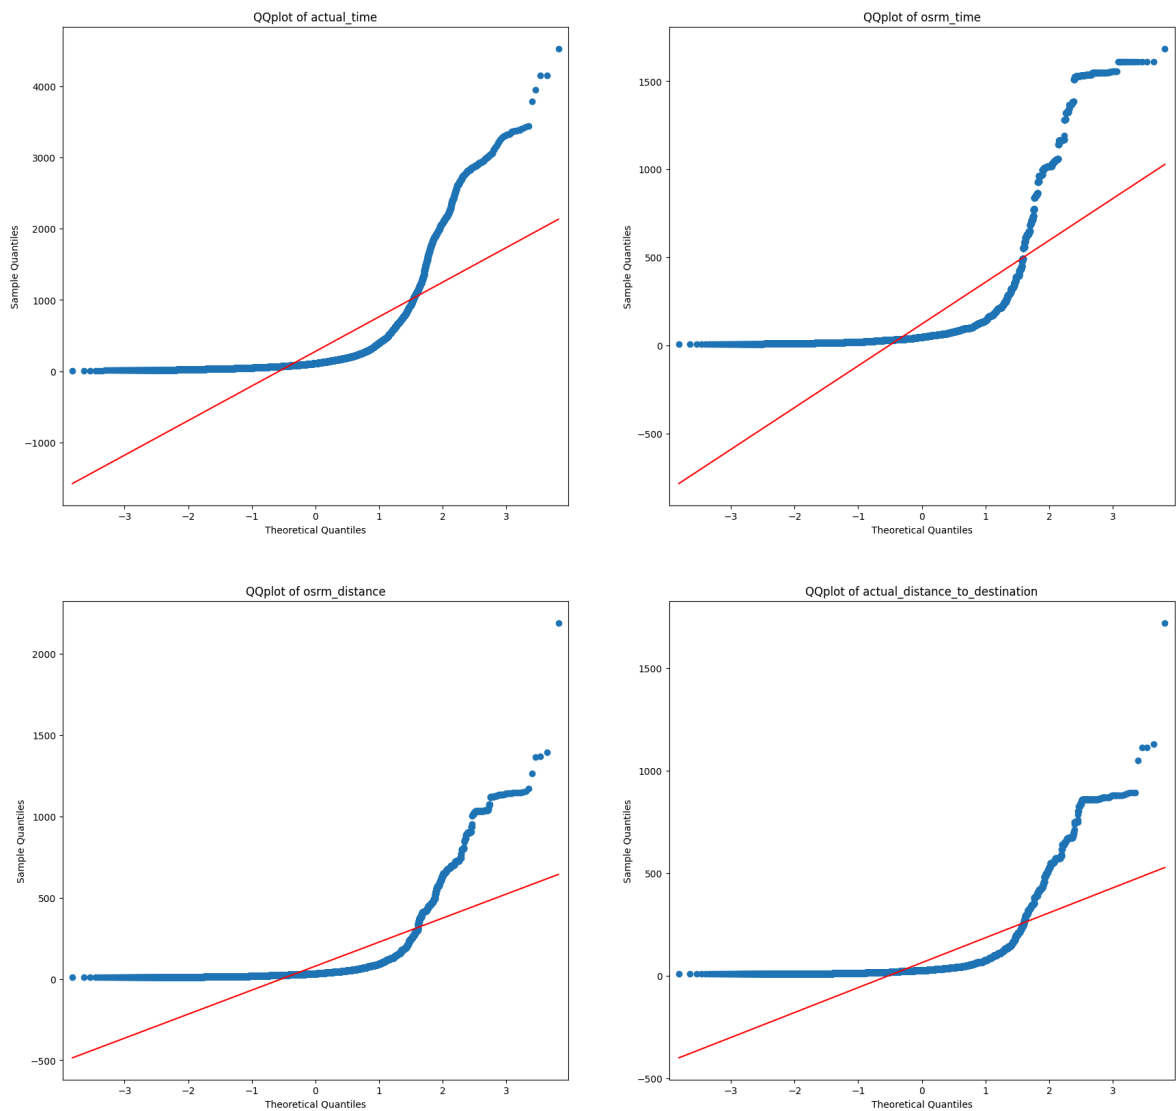
dtype: int64

```
In [17]: data_trip_uuid=data.groupby(['trip_uuid']).agg({'od_start_time':'max','od_end_time':'max','actual_distance':'sum','actual_time':'sum','osrm_time':'sum','osrm_distance':'sum','segment_actual_time':'sum','segment_osrm_time':'sum','segment_osrm_distance':'sum'}).reset_index()
```

```
In [18]: #Let us do a normality test using QQplot  
import scipy.stats as stats
```

```
In [20]: import statsmodels.api as sm
```

```
In [39]: fig, axis = plt.subplots(2, 2,figsize=(21, 20))  
  
sm.qqplot(data_trip_uuid['actual_time'],line='s',ax=axis[0,0])  
axis[0,0].set_title('QQplot of actual_time')  
  
sm.qqplot(data_trip_uuid['osrm_time'],line='s',ax=axis[0,1])  
axis[0,1].set_title('QQplot of osrm_time')  
  
sm.qqplot(data_trip_uuid['osrm_distance'],line='s',ax=axis[1,0])  
axis[1,0].set_title('QQplot of osrm_distance')  
  
sm.qqplot(data_trip_uuid['actual_distance_to_destination'],line='s',ax=axis[1,1])  
axis[1,1].set_title('QQplot of actual_distance_to_destination')  
  
plt.show()
```



- We can see that the data is not normally distributed so we can use Kruskal Wallis test for hypothesis testing

In [34]: `from scipy import stats`

```
In [40]: #segment_actual_time VS osrm_time
stat, p = stats.kruskal(data_trip_uuid['actual_time'], data_trip_uuid['osrm_time'])

print(f"Kruskal-Wallis H-statistic: {stat}")
print(f"P-value: {p}")

if p < 0.05:
    print("Significant difference between actual_time and osrm_time .")
else:
    print("No significant difference between actual_time and osrm_time.")
```

Kruskal-Wallis H-statistic: 4881.379995727001

P-value: 0.0

Significant difference between actual_time and osrm_time .

- By this we can tell that we need more accurate prediction of time through OSRM

```
In [41]: #actual_time VS segment_actual_time
stat, p = stats.kruskal(data_trip_uuid['actual_time'], data_trip_uuid['segment_ac

print(f"Kruskal-Wallis H-statistic: {stat}")
print(f"P-value: {p}")

if p < 0.05:
    print("Significant difference between actual_time and segment_actual_time .")
else:
    print("No significant difference between actual_time and segment_actual_time
```

Kruskal-Wallis H-statistic: 276.9587099963026

P-value: 3.454288758342234e-62

Significant difference between actual_time and segment_actual_time .

- Here also we can see there is a significant difference

```
In [42]: stat, p = stats.kruskal(data_trip_uuid['actual_distance_to_destination'], data_tr

print(f"Kruskal-Wallis H-statistic: {stat}")
print(f"P-value: {p}")

if p < 0.05:
    print("Significant difference between actual_distance_to_destination and osr
else:
    print("No significant difference between actual_distance_to_destination and
```

Kruskal-Wallis H-statistic: 774.5728435215805

P-value: 1.821772536232272e-170

Significant difference between actual_distance_to_destination and osrm_distance .

- I have tested three groups for comparing ORSM so the precision needs to be accurate .

Outliers

- Need to check if there are any outliers

```
In [50]: import matplotlib.pyplot as plt

fig, axis = plt.subplots(2, 2, figsize=(21, 20))

# Boxplot for actual_time
axis[0, 0].boxplot(data_trip_uuid['actual_time'].dropna())
axis[0, 0].set_title('Boxplot of actual_time')

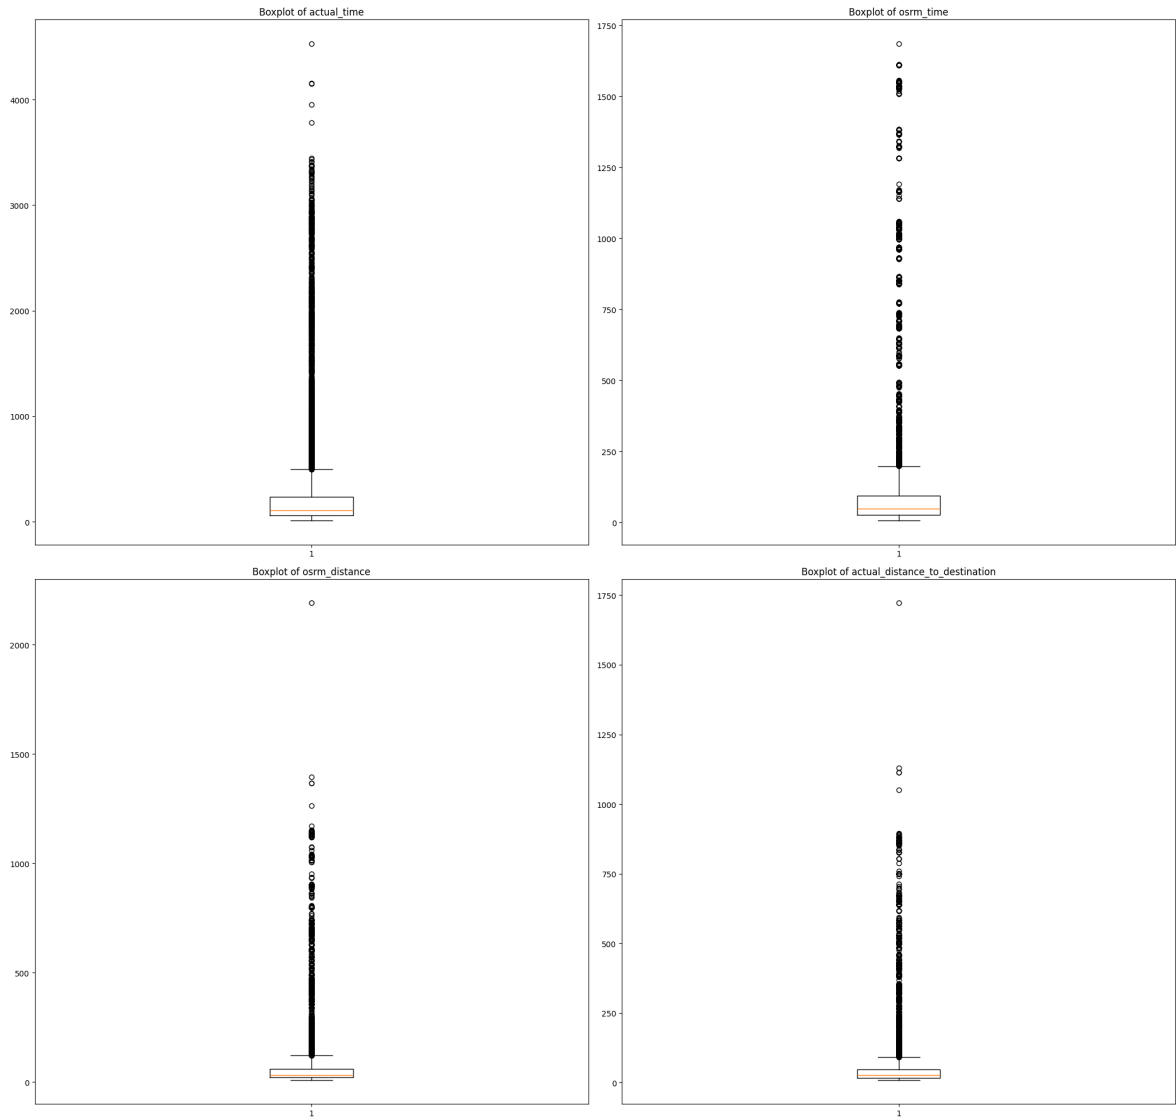
# Boxplot for osrm_time
axis[0, 1].boxplot(data_trip_uuid['osrm_time'].dropna())
axis[0, 1].set_title('Boxplot of osrm_time')

# Boxplot for osrm_distance
axis[1, 0].boxplot(data_trip_uuid['osrm_distance'].dropna())
```

```
axis[1, 0].set_title('Boxplot of osrm_distance')

# Boxplot for actual_distance_to_destination
axis[1, 1].boxplot(data_trip_uuid['actual_distance_to_destination'].dropna())
axis[1, 1].set_title('Boxplot of actual_distance_to_destination')

plt.tight_layout()
plt.show()
```



```
In [52]: columns = ['actual_time', 'osrm_time', 'osrm_distance', 'actual_distance_to_dest']

for col in columns:
    Q1 = data_trip_uuid[col].quantile(0.25)
    Q3 = data_trip_uuid[col].quantile(0.75)
    IQR = Q3 - Q1
    print(f"IQR for {col}: {IQR}")
```

IQR for actual_time: 174.0

IQR for osrm_time: 69.0

IQR for osrm_distance: 39.914138095238094

IQR for actual_distance_to_destination: 30.552448111226653

- There are outliers present in the data what we can do :

1. Remove Outliers

2. Cap or Winsorize Outliers
3. Log or Power Transform
4. Impute (if due to error or missing logic)
5. Label them

Business Insights

- Maharashtra has the highest number of orders both as a source and destination state. Tripura (as source) and Daman & Diu (as destination) have very low activity.
- most frequently used corridors
 1. IND562132AAA->IND560300AAA which is 151
 2. IND562132AAA ->IND560099AAB which is 127 These represent high-volume operational lanes.
- 1. Average OSRM Distance: 87.92
 2. Average OSRM Time: 121.69
 3. Average Actual Time: 273.54385543615297
- Tested if differences exist between groups (e.g., time prediction across routes). Found statistically significant differences → model needs to handle variability better.
- Found outliers in time and distance features. Suggested handling methods:
 1. Remove Outliers
 2. Cap or Winsorize Outliers
 3. Log or Power Transform
 4. Impute (if due to error or missing logic)
 5. Label them

Recommendations

- Focus marketing and expansion strategies on low-demand states to increase penetration.
- Optimize most frequently used corridors using real-time tracking, fleet management, and exception handling.
- Improve OSRM predictions using additional variables (traffic, weather, time of day).
- Also use robust ML models that don't assume normality.
- For outliers Use method that aligns with business logic.